Quiz: Sentiment Analysis Basics

12 multiple-choice questions. Choose one answer for each. Name: _____

1. What does a sentiment classifier predict?

A) The language of a review
B) Whether the opinion is positive or negative
C) The product category
D) The length of a review

2. Why do we keep the word 'not' during preprocessing?

A) It is rare
B) It speeds up training
C) It can flip the meaning of a phrase
D) It is a stop word that adds noise

3. What does TF-IDF do?

A) Weights words by how informative they are
B) Translates text
C) Removes spelling errors
D) Counts reviews

4. Which model is the baseline in this project?

A) DistilBERT
B) TF-IDF with Logistic Regression
C) k-means
D) A random guesser

5. Precision answers which question?

A) How many true positives were found
B) How many predicted positives were correct
C) How many reviews exist
D) How fast is the model

6. Recall answers which question?

A) How many real positives were found
B) How many words were removed
C) How many tokens fit
D) How many classes exist

7. F1 score is:

A) The mean of accuracy and loss
B) The harmonic mean of precision and recall
C) The number of errors
D) The training time

8. Why compare both models on the same test sample?

A) To make the comparison fair
B) To save disk space
C) Because the law requires it
D) To hide errors

9. Which review type is hardest for simple models?

A) Very short reviews with no verbs
B) Sarcastic reviews
C) Reviews in capital letters
D) Reviews with numbers

10. What is domain shift?

A) A model trained on one kind of text used on another
B) Moving files between folders
C) A change in random seed
D) A new GPU

11. Which claim is honest in the report?

A) The model is ready for production
B) Results are indicative on this dataset and sample
C) The model is never wrong
D) Accuracy proves fairness

12. Before trusting an AI-generated reference you should:

A) Copy it into the report
B) Open the source and confirm it exists
C) Ask the AI again
D) Ignore it